


# Serverless queues and workers — Designing Lift

Matthieu Napoli · Follow  
Published in Serverless Transformation · 4 min read · Apr 30, 2021

44

This article is part of a series on *Lift*, an open-source project that simplifies deploying serverless applications.




As introduced in previous articles, **Lift** is a **Serverless Framework plugin** that simplifies deploying serverless applications.

As we approach the first beta planned for May, let's talk about **queues and workers**:

- How to deploy production-ready queues and workers on AWS today.
- How we plan to simplify that with Lift.

## The naive approach



When going serverless, SQS + AWS Lambda is an iconic duo:

- SQS is a queue service that is entirely managed by AWS, scales automatically, and bills by the usage.
- AWS Lambda runs our code and is entirely managed by AWS, scales automatically, and bills by the usage.

The best part: Lambda integrates with SQS natively. We can write “queue workers” on Lambda without having to poll messages from SQS: instead, **our code (the worker) is automatically invoked when a new message (aka job) is pushed into SQS**.

To deploy queues and workers with the Serverless Framework, we need a bit of CloudFormation:

```
# serverless.yml
# ...

functions:
  worker:
    handler: worker.handler
    events:
      # Subscribe our function to the SQS queue
      - sqs:
          arn: !GetAtt ReportGenerationQueue.Arn

resources:
  Resources:
    # This creates an SQS queue
    ReportGenerationQueue:
      Type: AWS::SQS::Queue
```

So far so good! Or at least, it's doable.

## Problems with the naive approach

The biggest missing piece here is error handling.

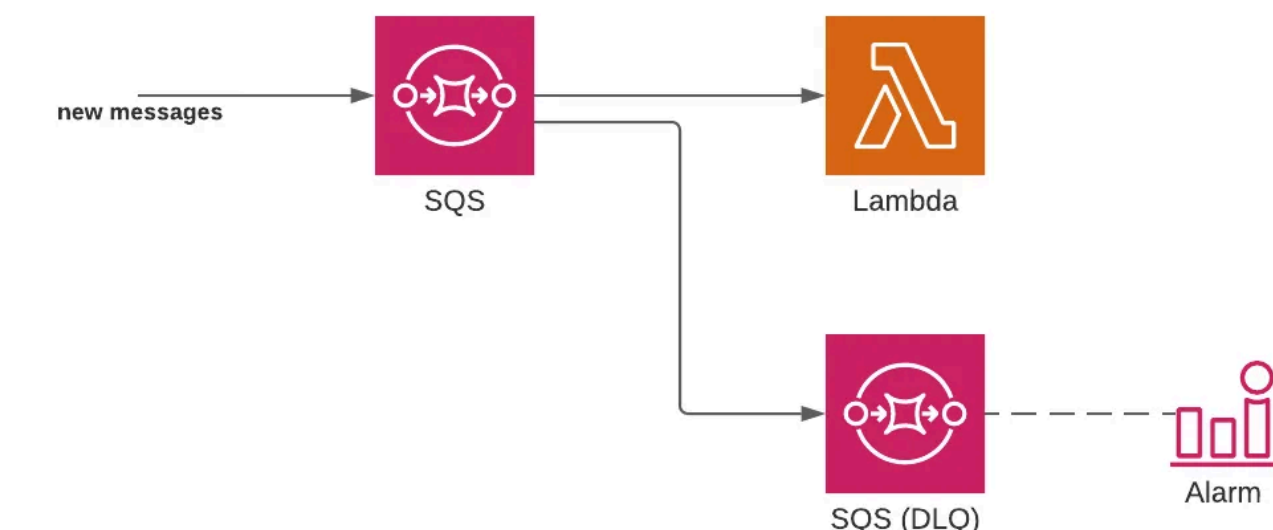
**By default, SQS retries failed jobs indefinitely.** Imagine your code has a bug (let's pretend that can happen): the job will be retried over and over for days, possibly costing a lot of money and wasting resources.

To deal with errors, we need to set up a “*dead letter queue*” (a queue that stores failed messages) and limit the number of retries.

On top of that, there are many details that needs to be configured, for example:

- We should configure the dead letter queue to store failed messages for 14 days (the maximum).
- We should set up an **alarm** that sends our team an email whenever there are failed messages in the dead letter queue.
- Many other settings need to be fine-tuned: the SQS “visibility timeout”, the Lambda max concurrency, message batching, etc.

## A production-ready approach



Here is a preview of a `serverless.yml` configuration that includes those best practices:

```
# serverless.yml
# ...

functions:
  worker:
    handler: worker.handler
    timeout: 10 # seconds
    # Scale to maximum 10 worker instances
    reservedConcurrency: 10
    events:
      - sqs:
          arn: !GetAtt ReportGenerationQueue.Arn
          # Process messages without batching
          batchSize: 1

resources:
  Resources:
    # Create an SQS queue
    ReportGenerationQueue:
      Type: AWS::SQS::Queue
      Properties:
        QueueName: !Sub '${AWS::StackName}-myqueue'
        # 6 times the lambda function's timeout
        VisibilityTimeout: 60
        # Messages will be stored (and retried) up to 2 days
        MessageRetentionPeriod: 172800
        RedrivePolicy:
          # Jobs will be retried 5 times
          maxReceiveCount: 5
          # Send failed messages to the dead letter queue
          deadLetterTargetArn: !GetAtt DeadLetterQueue.Arn

    # Failed messages will be sent to this queue
    DeadLetterQueue:
      Type: AWS::SQS::Queue
      Properties:
```

```
# Store messages up to 14 days (the max)
# we want time to debug/analyse
MessageRetentionPeriod: 1209600
```

If we wanted to add an alarm for failed messages, we would need 30 more lines of YAML.

Needless to say:

- this is hard to figure out,
- this is tedious to implement.

We want to make that simpler with Lift.

### Serverless queues and workers with Lift

We are currently working on a “Queues” component that can be deployed via `serverless.yml`:

```
# serverless.yml
# ...

queues:
  report-generation:
    # Our Lambda worker function is defined in the "queue" component:
    worker:
      handler: worker.handler
      # Scale to maximum 10 worker instances
      reservedConcurrency: 10
```

As we can see in the example above, there is a new `queues` section that lets us define queues (and their workers). In that instance, we define a `report-generation` queue and its worker function.

On `serverless deploy`, Lift will create:

- A `report-generation` SQS queue configured following best practices.
- A `worker` Lambda function, that will be automatically subscribed to the SQS queue.
- A `report-generation-dlq` SQS “dead letter queue” that will receive failed messages.

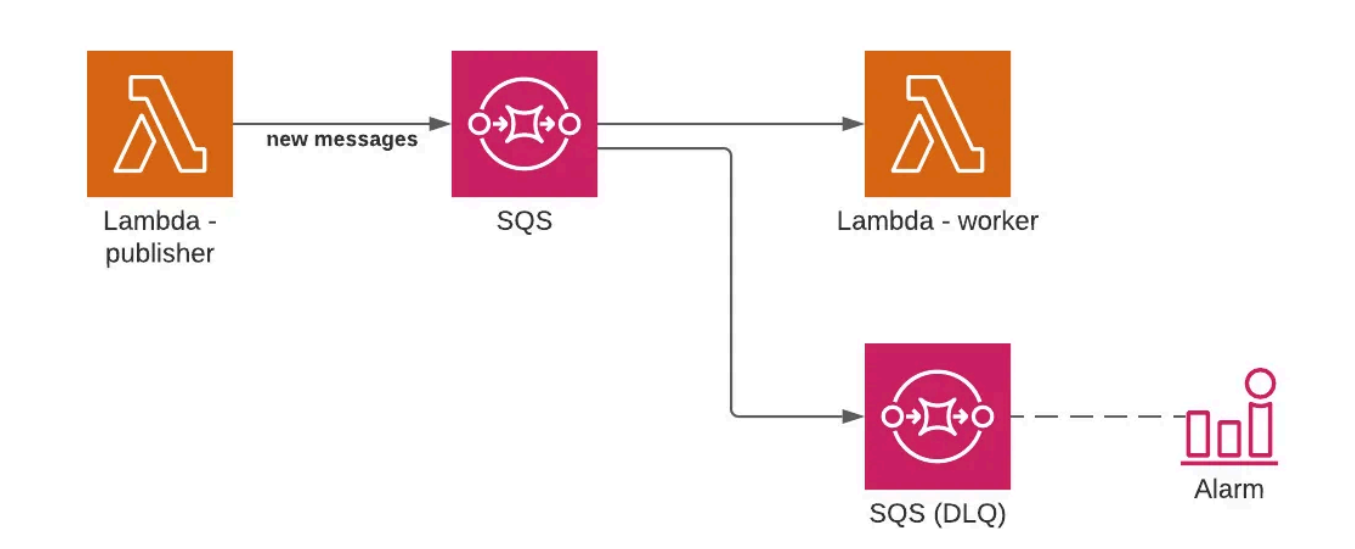
By default, messages are configured to be retried up to 3 times.

We are considering several settings on the component, including setting an email alert on failed messages:

```
# serverless.yml
# ...

queues:
  report-generation:
    alarm: alerting@my-company.com
    worker:
      ...
```

On top of that, Lift also simplifies **pushing messages** into the SQS queue:



- Lambda functions defined in `serverless.yml` will automatically get IAM **permissions** to send messages to the SQS queue.
- To simplify referencing the SQS queue, Lift introduces a **new variable syntax**: `${queues:my-queue.queueUrl}`

Here is a complete example with a *publisher* Lambda function:

```
# serverless.yml
# ...

functions:
  publisher:
    handler: publisher.handler
    environment:
      QUEUE_URL: ${queues:report-generation.queueUrl}


queues:
  report-generation:
    worker:
      handler: worker.handler
```

Deploying queues, [static websites](#) and [file storage](#) is a small part of what we are working on with Lift.

To follow Lift’s development, [star or watch Lift on GitHub](#).

We are also looking for feedback on the *Queues* feature: [get involved in the GitHub discussion](#).

Serverless   AWS   Lift   Queue   Asynchronous



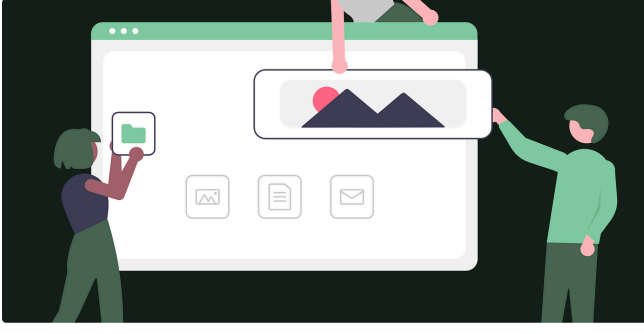
**Written by Matthieu Napoli**

234 Followers · Writer for Serverless Transformation

Make serverless more accessible

Follow

More from Matthieu Napoli and Serverless Transformation

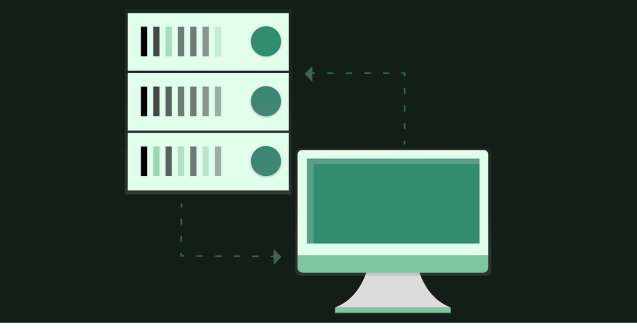


Matthieu Napoli in Serverless Transformation

**Static websites on AWS — Designing Lift**

This article is part of a series on Lift, an open-source project that simplifies deploying...

3 min read · Apr 16, 2021

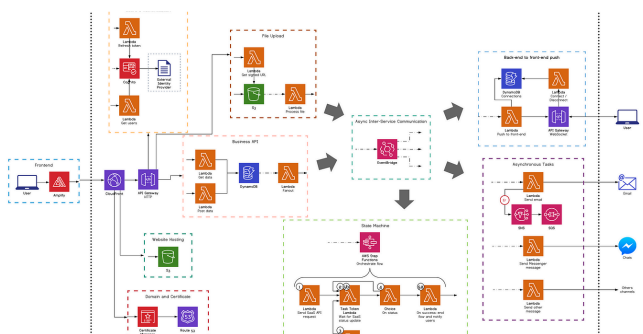


Xavier Lefèvre in Serverless Transformation

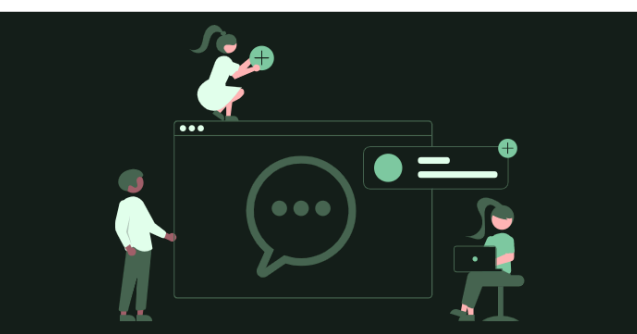
**Asynchronous client interaction in AWS Serverless: Polling,...**

Event-driven Serverless often requires async update to the client. There are several ways ...

6 min read · Apr 18, 2020



Xavier Lefèvre in Serverless Transformation



Sarah Hamilton in Serverless Transformation